

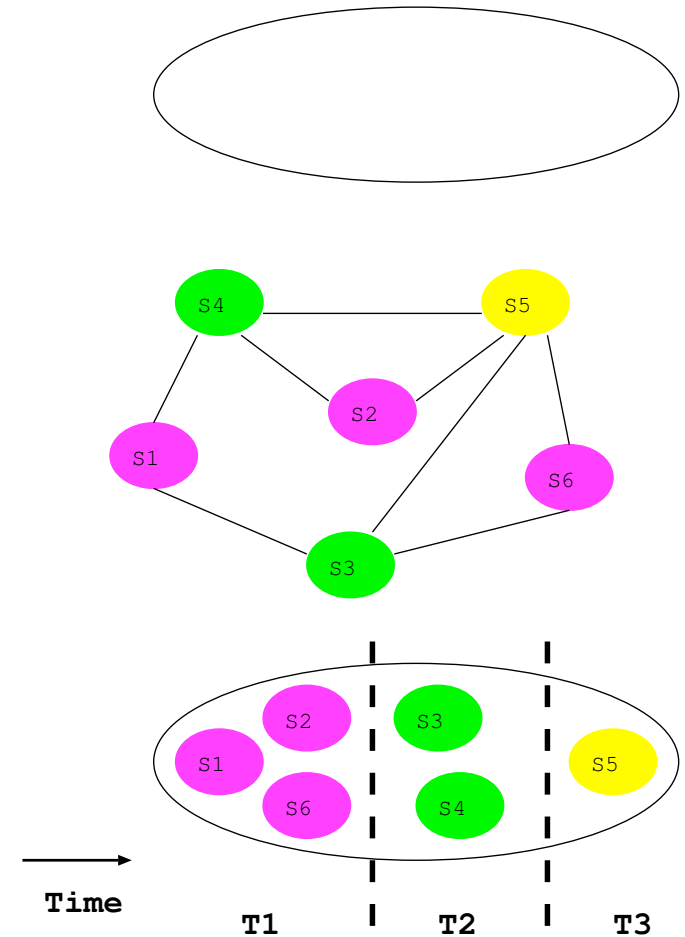
Parallelizing Graph Algorithms

- Challenges:
 - Runtime dominated by memory latency than processor speed
 - Little work is done while visiting a vertex or an edge
 - Little computation to hide memory access cost
 - Access patterns determined only at runtime
 - Prefetching techniques inapplicable
 - There is poor data locality
 - Difficult to obtain good memory system performance
- For these reasons, parallel performance
 - on distributed memory machines is often poor
 - on shared memory machines is often better

We consider here *graph coloring* as an example of a graph algorithm to parallelize on shared memory machines

Graph coloring

- *Graph coloring* is an assignment of colors (positive integers) to the vertices of a graph such that adjacent vertices get different colors
- The objective is to find a coloring with the *least* number of colors
- Examples of *applications*:
 - Concurrency discovery in parallel computing (illustrated in the figure to the right)
 - Sparse derivative computation
 - Frequency assignment
 - Register allocation, etc



A greedy algorithm for coloring

- Graph coloring is NP-hard to solve optimally (and even to approximate)
- The following Greedy algorithm gives very good solution in practice

Algorithm 1 Sequential greedy coloring.

```
1: procedure GREEDY( $G(V, E)$ )
2:   for each  $v \in V$  do
3:     for each  $w \in \text{adj}(v)$  do
4:        $\text{forbiddenColors}[\text{color}[w]] \leftarrow v$        $\triangleright$  mark color of  $w$  as forbidden to  $v$ 
5:        $\text{color}[v] \leftarrow \min\{c > 0 : \text{forbiddenColors}[c] \neq v\}$        $\triangleright c$  is the smallest
        permissible color to  $v$ 
```

color is a vertex-indexed array that stores the color of each vertex

forbiddenColors is a color-indexed array used to mark impermissible colors to a vertex

Complexity of GREEDY: $O(|E|)$ (thanks to the way the array *forbiddenColors* is used)

Parallelizing Greedy Coloring

- Desired goal: parallelize GREEDY such that
 - Parallel runtime is roughly $O(|E|/p)$ when p processors (threads) are used
 - Number of colors used is nearly the same as in the serial case
- Difficult to achieve since GREEDY is inherently sequential
- Challenge: come up with a way to create concurrency in a nontrivial way

A potentially “generic” parallelization technique

- “Standard” Partitioning
 - Break up the given problem into p independent subproblems of almost equal sizes
 - Solve the p subproblems concurrently

Main work lies in the decomposition step which is often no easier than solving the original problem

- “Relaxed” Partitioning
 - Break up the problem into p , not necessarily entirely independent, subproblems of almost equal sizes
 - Solve the p subproblems concurrently
 - Detect inconsistencies in the solutions concurrently
 - Resolve any inconsistencies

Can be used potentially successfully if the resolution in the fourth step involves only local adjustments

“Relaxed Partitioning” applied towards parallelizing Greedy coloring

- Speculation and Iteration:
 - Color as many vertices as possible concurrently, tentatively tolerating potential conflicts, detect and resolve conflicts afterwards (iteratively)


Parallel Coloring on Shared Memory Platforms (using Speculation and Iteration)

Algorithm 2 Iterative parallel greedy coloring.

```
1: procedure ITERATIVE( $G(V, E)$ )
2:    $U \leftarrow V$  ▷  $U$  is the current set of vertices to be colored
3:   while  $U \neq \emptyset$  do
4:     for each  $v \in U$  in parallel do ▷ Phase 1: tentative coloring
5:       for each  $w \in \text{adj}(v)$  do
6:         mark  $\text{color}[w]$  as forbidden to  $v$ 
7:         Pick the smallest permissible color  $c$  for vertex  $v$ 
8:        $R \leftarrow \emptyset$  ▷  $R$  is a set of vertices to be recolored
9:       for each  $v \in U$  in parallel do ▷ Phase 2: conflict detection
10:        for each  $w \in \text{adj}(v)$  do
11:          if  $\text{color}[v] = \text{color}[w]$  and  $v > w$  then
12:             $R \leftarrow R \cup \{v\}$ 
13:        $U \leftarrow R$ 
```

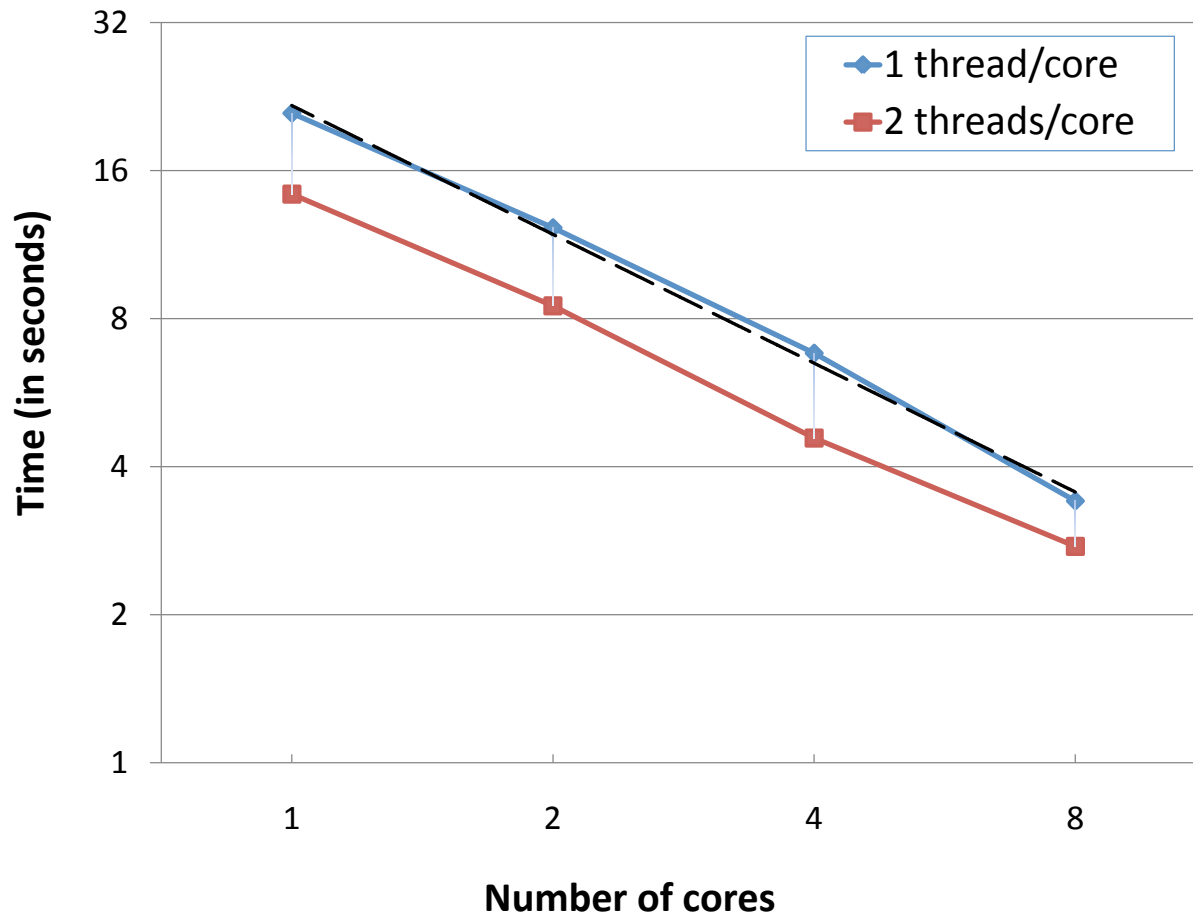
Lines 4 and 9 can be parallelized using the OpenMP directive
#pragma omp parallel for

Sample experimental results of Algorithm 2 (Iterative) on Nehalem: I



Graph (RMAT-G): 16.7M vertices; 133.1M edges;
Degree (avg=16, Max= 1,278, variance=416)

Sample experimental results of Algorithm 2 (Iterative) on Nehalem: II



Graph (RMAT-B): 16.7M vertices; 133.7M edges;
Degree (avg=16, Max= 38,143, variance=8,086)